

Information Retrieval & Social Web

CS 525/DS 595

Worcester Polytechnic Institute

Department of Computer Science

Instructor: Prof. Kyumin Lee

Project

- 3 or 4 person team
- Dates:
 - [7%] Project Proposal: March 18 by 11:59pm
 - [8%] Project website: April 24 by 11:59pm
 - [16%] Final project presentation: April 25 in-class
- https://canvas.wpi.edu/courses/7874/discussion_topics/31731

Previous Class...

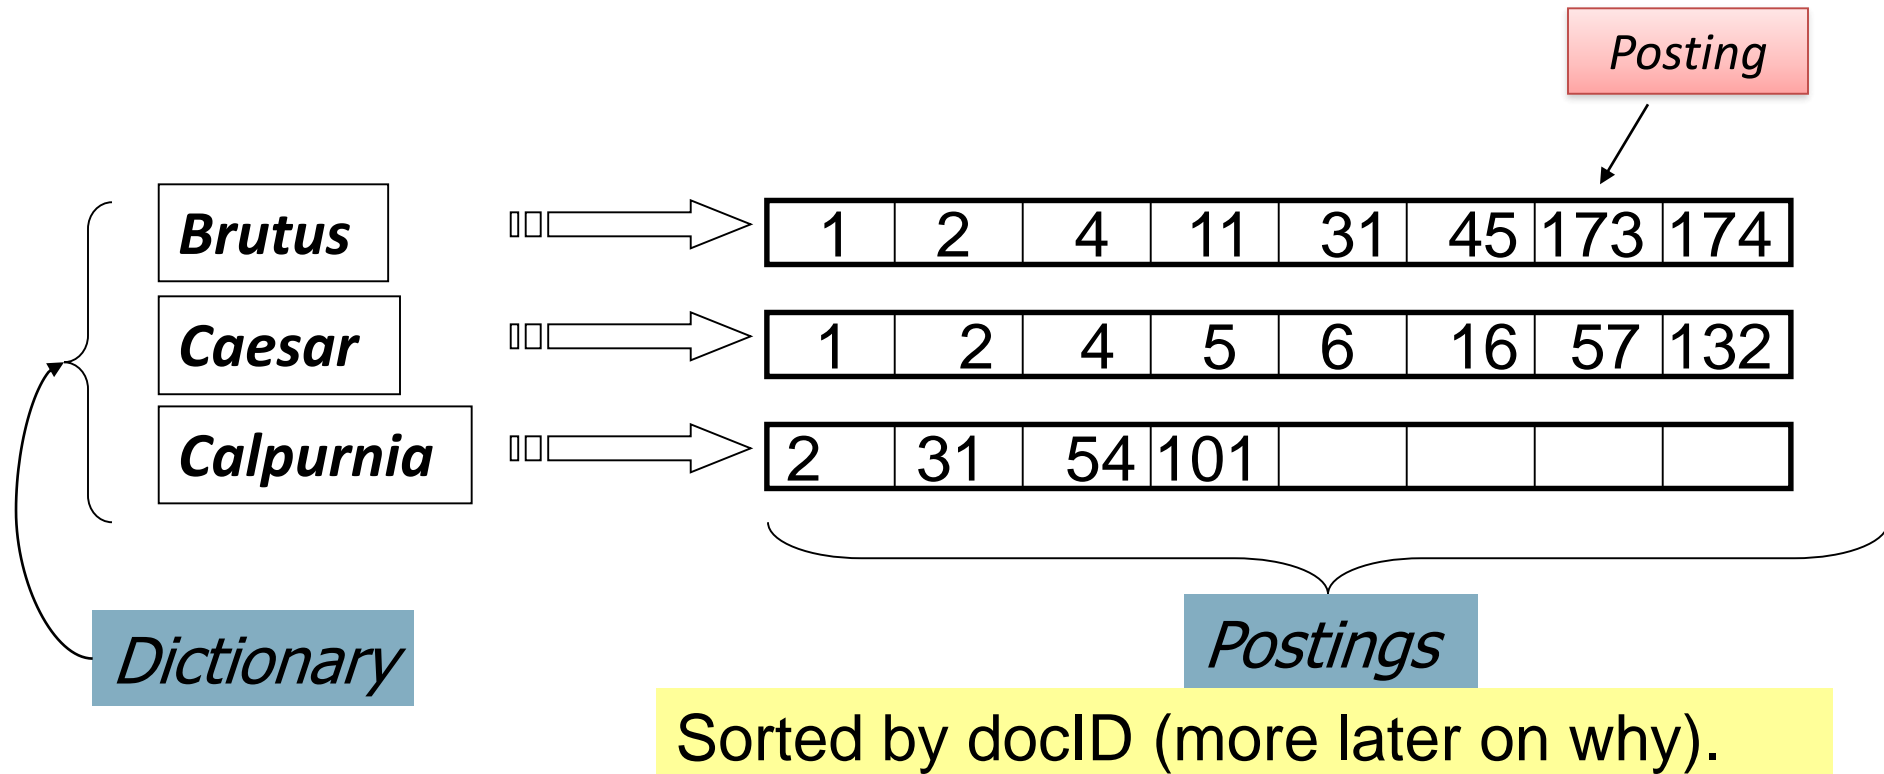
Boolean Retrieval
Model

Previous Class...

Boolean Retrieval
Model

Inverted index

Inverted index



Boolean queries: More general merges

- Exercise: Adapt the merge for the queries:

Brutus AND NOT Caesar

Brutus OR NOT Caesar

Can we still run through the merge in time $O(x+y)$?

What can we achieve?

Exercise Solution

- Brutus AND NOT Caesar
 - Time is $O(x+y)$. Instead of collecting documents that occur in both postings lists, collect those that occur in the first one and not in the second
- Brutus OR NOT Caesar
 - Time is $O(N)$ (where N is the total number of documents in the collection) assuming we need to return a complete list of all documents satisfying the query. This is because the length of the result list is only bounded by N , not by the length of the postings lists.

Merging

What about an arbitrary Boolean formula?

(Brutus OR Caesar) AND NOT

(Antony OR Cleopatra)


- Can we always merge in “linear” time?
 - Linear in what?
- Can we do better?

Solution

- We can always intersect in $O(qN)$ where q is the number of query terms and N the number of documents, so the intersection time is linear in the number of documents and query terms. Since the tightest bound for the size of the result list is N , the number of documents, one cannot do better than $O(N)$.
- But... still we can reduce computation time even though time complexity is still $O(N)$. How?

Query optimization

- Consider a query that is an *AND* of n terms.
- For each of the n terms, get its postings, then *AND* them together.
- What is the best order for query processing?




Query: **Brutus AND Calpurnia AND Caesar**

Query optimization example

- Process in order of increasing freq:
 - *start with smallest set, then keep cutting further.*

This is why we kept
document freq. in dictionary



Execute the query as (***Calpurnia AND Brutus***) ***AND Caesar***.

More general optimization

- e.g., (*madding OR crowd*) AND (*ignoble OR strife*)
- Get doc. freq.'s for all terms.
- Estimate the size of each *OR* by the sum of its doc. freq.'s (conservative).
- Process in increasing order of *OR* sizes.

Exercise

- Recommend a query processing order for

*(tangerine OR trees) AND
(marmalade OR skies) AND
(kaleidoscope OR eyes)*

| Term | Freq |
|--------------|--------|
| eyes | 213312 |
| kaleidoscope | 87009 |
| marmalade | 107913 |
| skies | 271658 |
| tangerine | 46653 |
| trees | 316812 |

Exercise Solution

- Using the conservative estimate of the length of unioned postings lists, the recommended order is: (kaleidoskope OR eyes) (300,321) AND (tangerine OR trees) (363,465) AND (marmalade OR skies) (379,571)

What's ahead in IR?

Beyond term search

- What about phrases?
 - *Stanford University*
- Proximity: Find ***Gates NEAR Microsoft.***
 - Need index to capture position information in docs.
- Zones in documents: Find documents with
(*author = Ullman*) AND (text contains ***automata***).

Evidence accumulation

- 1 vs. 0 occurrence of a search term
 - 2 vs. 1 occurrence
 - 3 vs. 2 occurrences, etc.
 - Usually more seems better
- Need term frequency information in docs

Ranking search results

- Boolean queries give inclusion or exclusion of docs.
- Often we want to rank/group results
 - Need to measure proximity from query to each doc.
 - Need to decide whether docs presented to user are singletons, or a group of docs covering various aspects of the query.

Clustering, classification and ranking

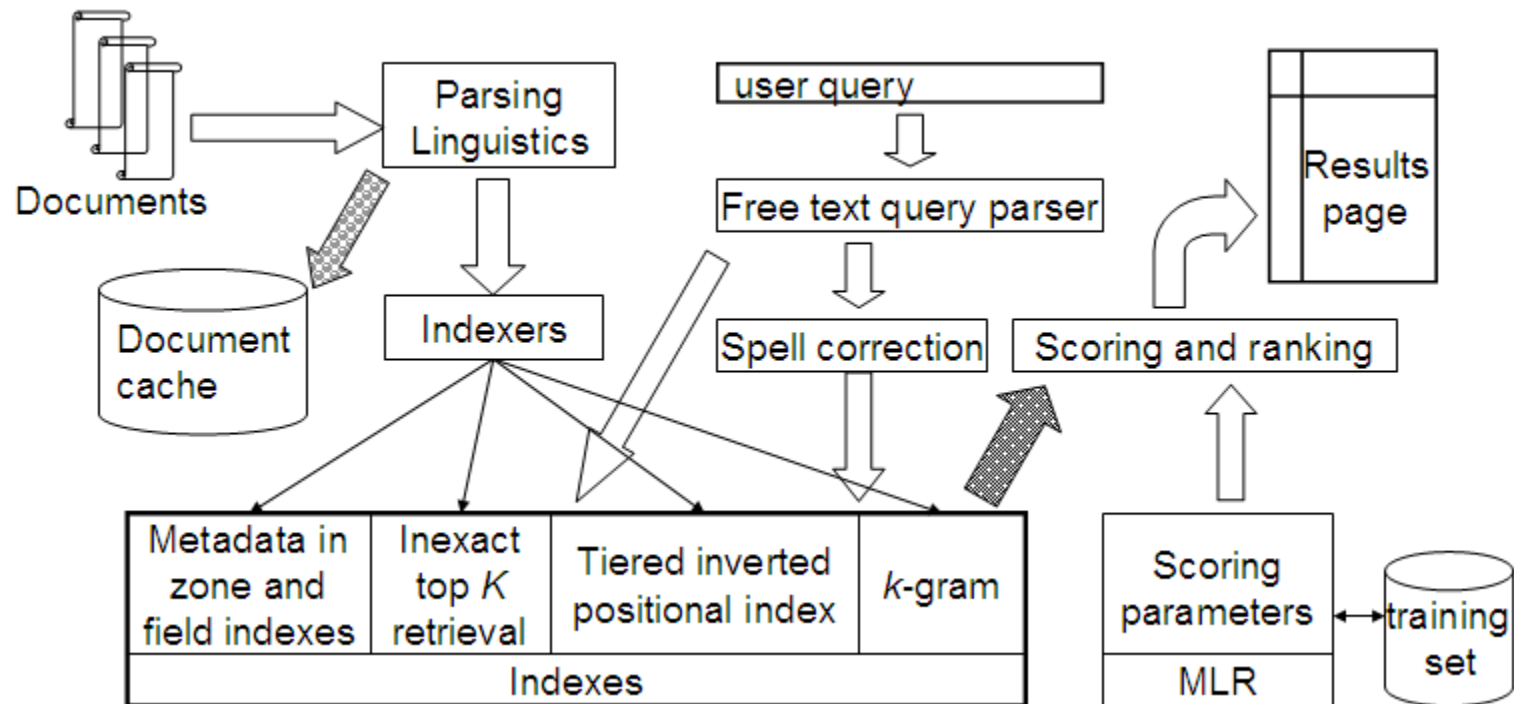
- **Clustering:** Given a set of docs, group them into clusters based on their contents.
- **Classification:** Given a set of topics, plus a new doc D , decide which topic(s) D belongs to.
- **Ranking:** Can we learn how to best order a set of documents, e.g., a set of search results

Exercise: Build Inverted Index and Run Boolean Retrieval

- Doc1: The winning ticket in Florida was sold at a Publix Supermarket.
- Doc2: The winning numbers were 08, 27, 34, 04 and 19, and the Powerball was 10.
- Doc3: With three winning tickets, the lump sum will be \$187.2 million.

- Doc1 AND Doc2 AND Doc3
- Doc1 OR Doc2

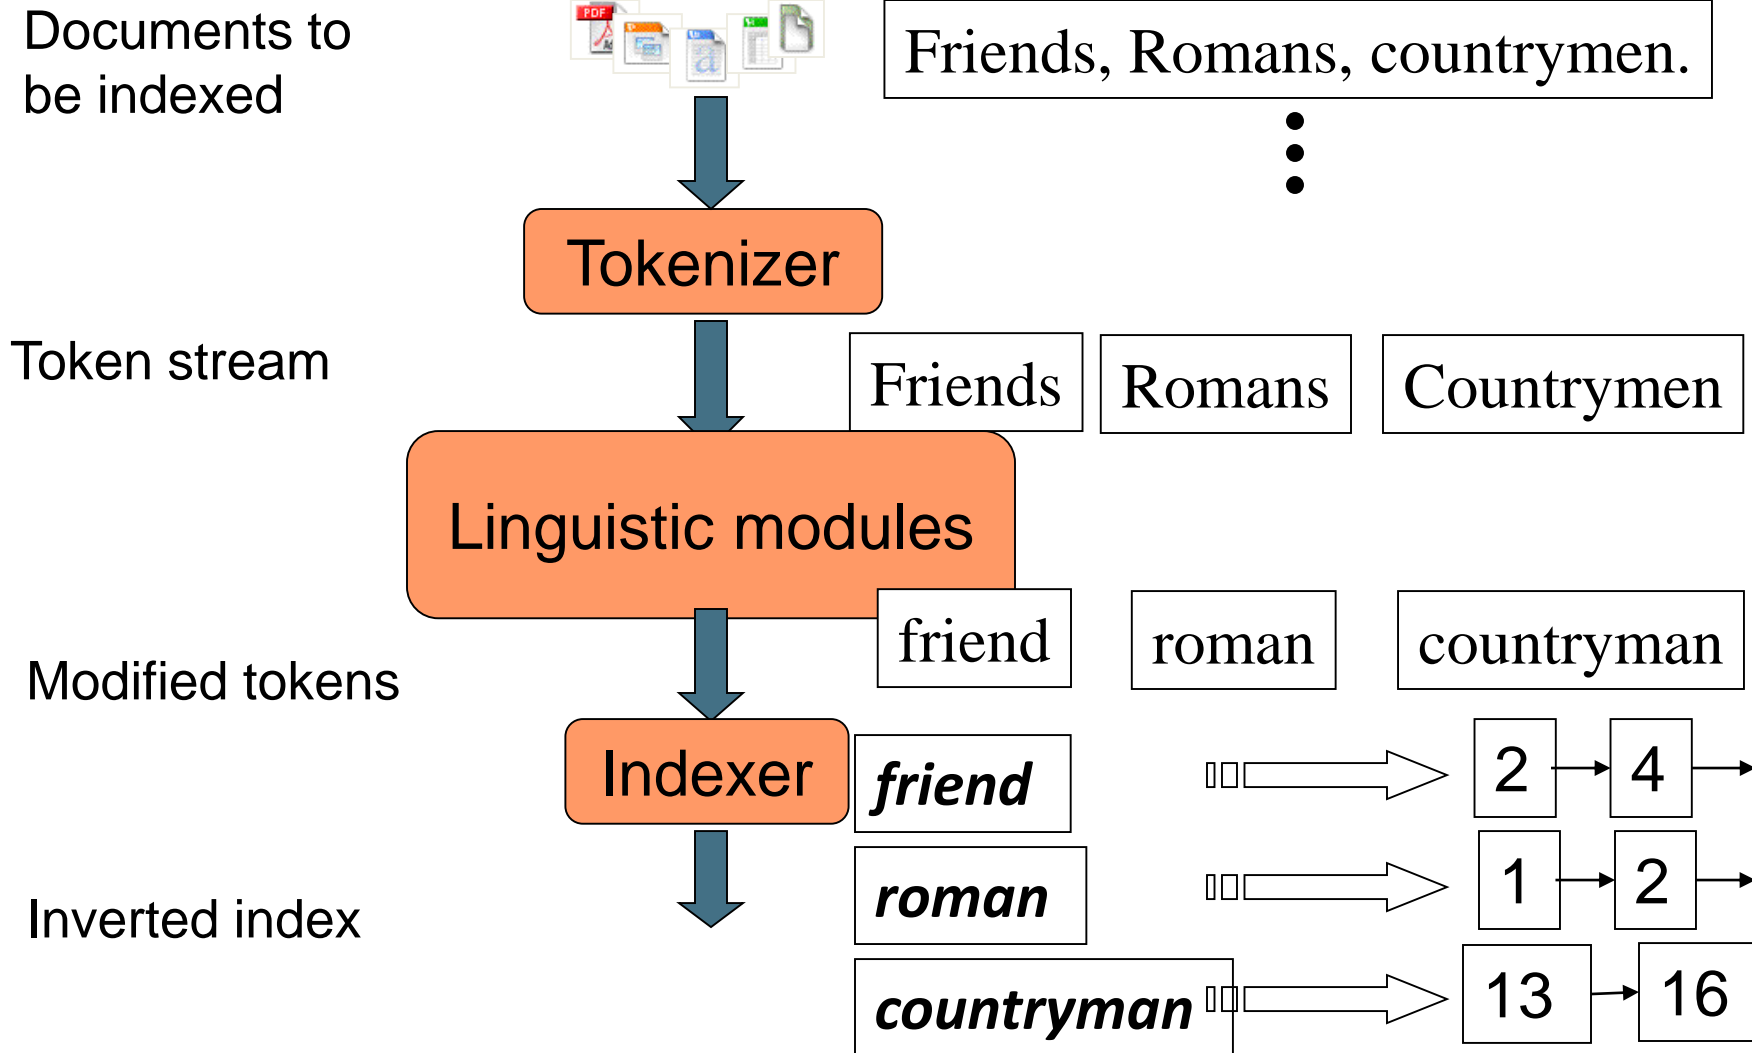
What's next ...



Our assumptions so far

- We know what a document is
- We know what a term is
 - In reality, it can be complex
- So... We'll look at how we define and process the vocabulary of terms in a collection

Recall the basic indexing pipeline



Initial stages of text processing

- Tokenization
 - Cut character sequence into word tokens
 - Deal with ***“John’s”, a state-of-the-art solution***
- Normalization
 - Map text and query term to same form
 - You want **U.S.A.** and **USA** to match
- Stemming
 - We may wish different forms of a root to match
 - ***authorize, authorization***
- Stop words
 - We may omit very common words (or not)
 - ***the, a, to, of***

Parsing a document

- What format is it in?
 - pdf/word/excel/html?
- What language is it in?
- What character set is in use?
 - (CP1252, UTF-8, ...)

Each of these is a classification problem, which we will study later in the course.

But these tasks are often done heuristically ...

Complications: Format/language

- Documents being indexed can include docs from many different languages
 - A single index may contain terms from many languages.
- Sometimes a document or its components can contain multiple languages/formats
 - French email with a German pdf attachment.
- What is a unit document?
 - A file?
 - An email? (Perhaps one of many in a single mbox file)
 - An email with 5 attachments?
 - A group of files (PPT or LaTeX in HTML)

Tokenization

- Input: “*Friends, Romans and Countrymen*”
- Output: Tokens
 - *Friends*
 - *Romans*
 - *Countrymen*
- A **token** is an instance of a sequence of characters
- Each such token is now a candidate for an index entry, after further processing
 - Described below
- But what are valid tokens to emit?

Why tokenization is difficult -- even in English

- Example: *Mr. O'Neill thinks that the boys' stories about Chile's capital aren't amusing.*
- **Tokenize this sentence**

One word or two? (or several)

- Hewlett-Packard
- State-of-the-art
- co-education
- the hold-him-back-and-drag-him-away maneuver
- data base
- San Francisco
- Los Angeles-based company
- cheap San Francisco-Los Angeles fares
- York University vs. New York University

Numbers

- 3/12/91
- 12/3/91
- Mar 12, 1991
- B-52
- 100.2.86.144
- (800) 234-2333
- 800.234.2333

Chinese: No whitespace

莎拉波娃现在居住在美国东南部的佛罗里达。今年4月9日，莎拉波娃在美国第一大城市纽约度过了18岁生日。生日派对上，莎拉波娃露出了甜美的微笑。

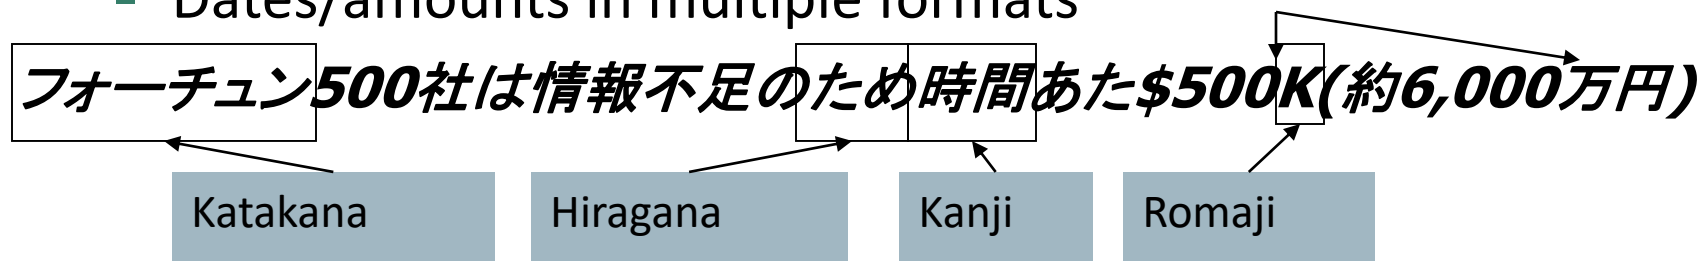
Ambiguous segmentation in Chinese

和尚

- Can be treated as one word meaning “monk” or as two words meaning “and” and “still”

Tokenization: language issues

- Chinese and Japanese have no spaces between words:
 - 莎拉波娃现在居住在美国东南部的佛罗里达。
 - Not always guaranteed a unique tokenization
- Further complicated in Japanese, with multiple alphabets intermingled
 - Dates/amounts in multiple formats



End-user can express query entirely in hiragana!

Language issues in French

- ***L'ensemble*** → one token or two?
 - ***L ? L' ? Le ?***
 - Want ***l'ensemble*** to match with ***un ensemble***

Bidirectionality in Arabic

- Arabic (or Hebrew) is basically written right to left, but with certain items like numbers written left to right
 - Words are separated, but letter forms within a word form complex ligatures
- استقلت الجزائر في سنة 1962 بعد 132 عاما من الاحتلال الفرنسي.
- ← → ← → ← start
- ‘Algeria achieved its independence in 1962 after 132 years of French occupation.’

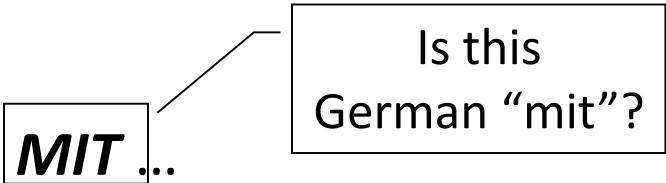
Normalization

- Need to “normalize” words in indexed text as well as query words into the same form
 - We want to match **U.S.A.** and **USA**
- We most commonly implicitly define **equivalence classes** of terms
 - e.g., deleting periods to form a term
- Alternative is to do asymmetric expansion:
 - Enter: *window* Search: *window, windows*
 - Enter: *windows* Search: *Windows, windows*
 - Enter: *Windows* Search: *Windows*
- Potentially more powerful, but less efficient

Normalization: other languages

- Accents: e.g., French *résumé* vs. *resume*.
- Most important criterion:
 - How are your users like to write their queries for these words?
- Even in languages that standardly have accents, users often may not type them
- German: *Tuebingen* vs. *Tübingen*
 - Should be equivalent

Normalization: other languages

- Need to “normalize” indexed text as well as query terms into the same form
 - 7月30日 vs. 7/30
- Character-level alphabet detection and conversion
 - Tokenization not separable from this.
 - Sometimes ambiguous:
Morgen will ich in MIT ..

Is this
German “mit”?

Case folding

- Reduce all letters to lower case
 - exception: upper case in mid-sentence?
 - e.g., General Motors
 - Fed vs. fed
 - SAIL vs. sail
- Often best to lower case everything, since users will use lowercase regardless of 'correct' capitalization...

Stop words

- With a stop list, you exclude from the dictionary entirely the commonest words. Intuition:
 - They have little semantic content: *the, a, and, to, be*
 - There are a lot of them: ~30% of postings for top 30 words
- But the trend is away from doing this:
 - Good compression techniques means the space for including stop words in a system is very small
 - Good query optimization techniques mean you pay little at query time for including stop words.
 - You need them for:
 - Phrase queries: “King of Denmark”
 - Various song titles, etc.: “Let it be”, “To be or not to be”
 - “Relational” queries: “flights to London”

Lemmatization

- Reduce inflectional/variant forms to base form
- Example: *am, are, is* → *be*
- Example: *car, cars, car's, cars'* → *car*
- Example: the boy's cars are different colors → the boy car be different color
- Lemmatization implies doing “proper” reduction to dictionary headword form (the lemma).

Stemming

- Reduce terms to their “roots” before indexing
- “Stemming” suggests crude affix chopping
- language dependent
- Example: ***automate(s), automatic, automation*** all reduced to ***automat***.

Porter Stemming Algorithm

- Most common algorithm for stemming English
- Results suggest that it is at least as good as other stemming options
- Contains 5 phases of reductions
- Phases are applied sequentially
- Each phase consists of a set of commands.
 - Sample command: Delete final element if what remains is longer than 1 character
 - replacement → replac
 - cement → cement

Porter stemmer:

A few rules

Rule

SSSES → SS

IES → I

SS → SS

S →

Example

caresses → caress

ponies → poni

caress → caress

cats → cat

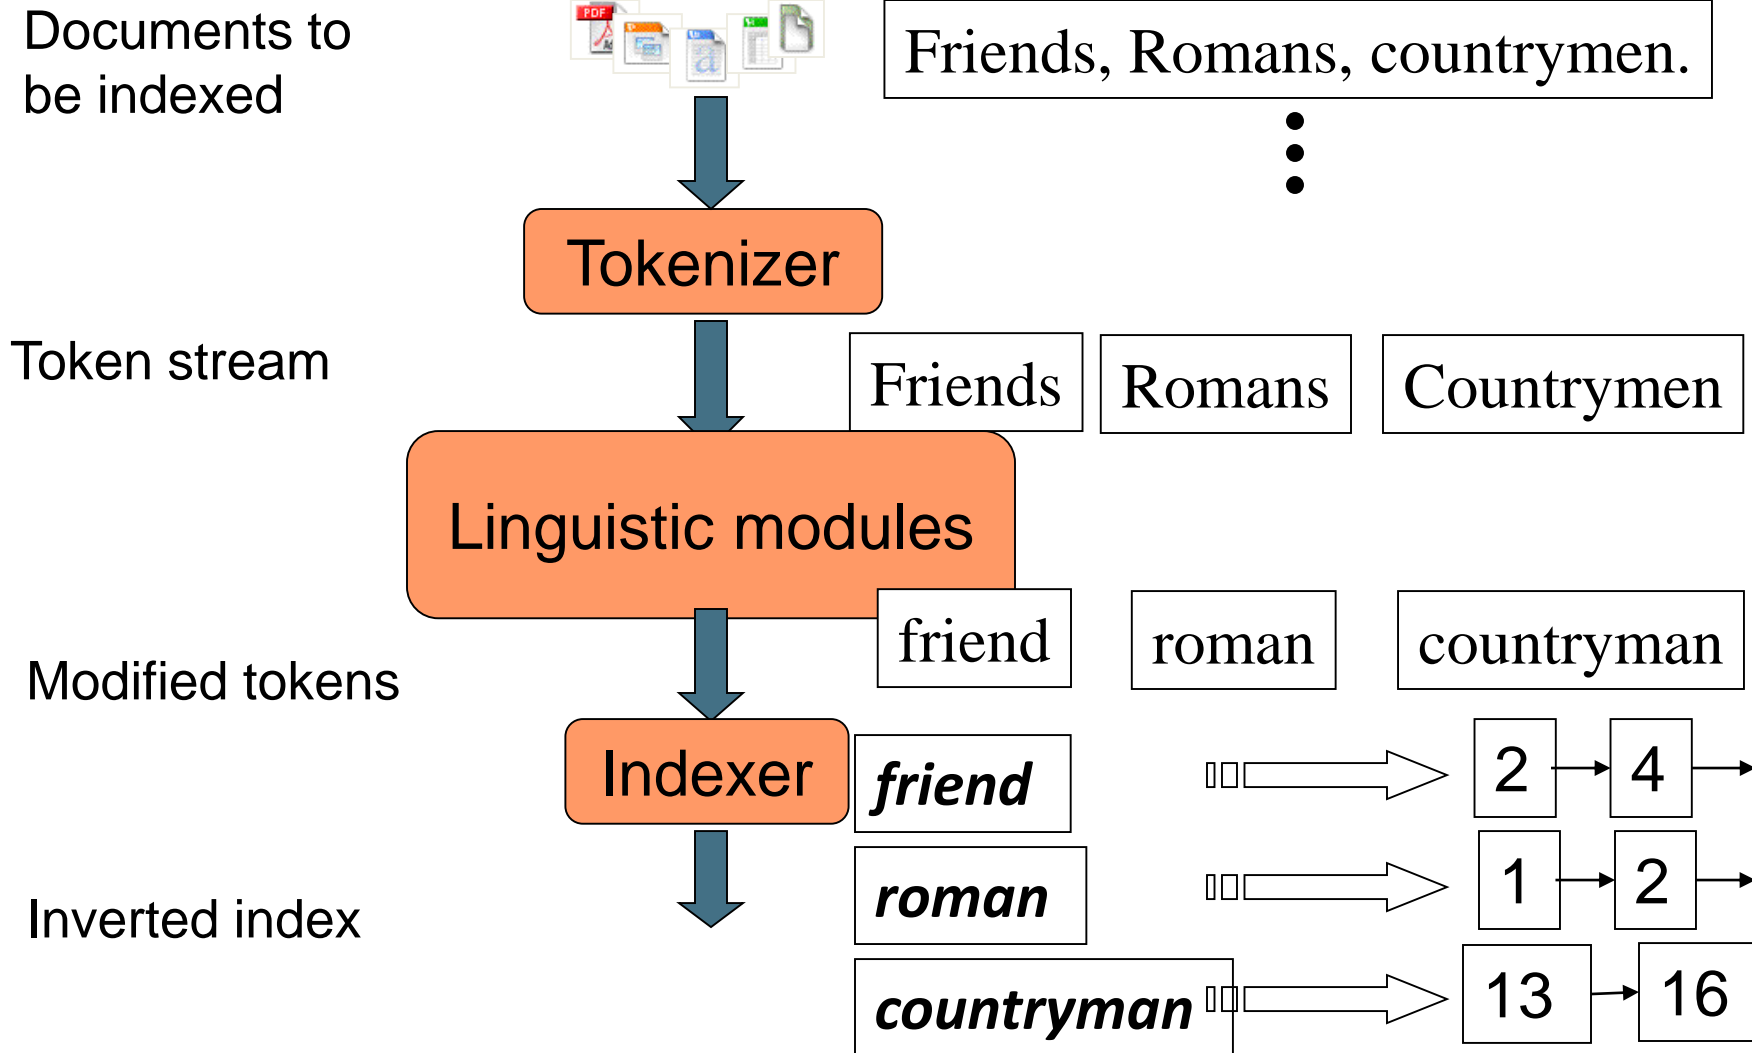
Three stemmers: A comparison

- **Sample text:** Such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation
- **Porter stemmer:** such an analysis can reveal features that are not easily visible from the variation in the individual gene and can lead to a picture of expression that is more biologically transparent and accessible to interpretation
- **Lovins stemmer:** such an analysis can reveal features that are not easily visible from the variation in the individual gene and can lead to a picture of expression that is more biologically transparent and accessible to interpretation
- **Paice stemmer:** such an analysis can reveal features that are not easily visible from the variation in the individual gene and can lead to a picture of expression that is more biologically transparent and accessible to interpretation

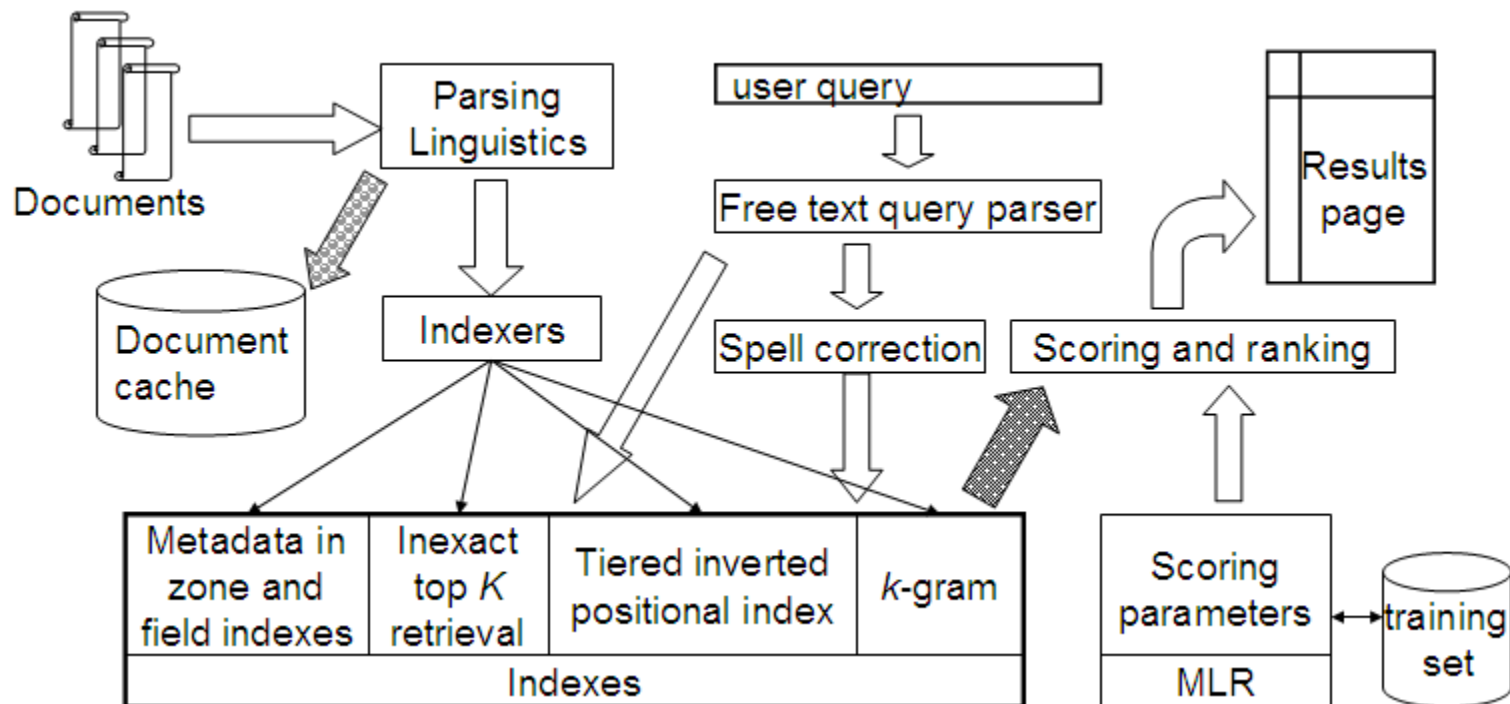
Does stemming improve effectiveness?

- In general, stemming increases effectiveness for some queries, and decreases effectiveness for others.
- Porter Stemmer equivalence class **oper** contains all of **operate operating operates operation operative operatives operational**.
- Queries where stemming hurts: “operational AND research”, “operating AND system”, “operative AND dentistry”

Recall the basic indexing pipeline



Big Picture



HW1

<https://canvas.wpi.edu/courses/7874/assignments/477>
26


Next...

- Need a better index than simple <term: docs>
- How can we improve on our basic index?
 - **Skip pointers**: faster postings merges
 - **Positional index**: Phrase queries and Proximity queries
 - **Permuterm index**: Wildcard queries
 - **k-gram index**: Wildcard queries and spell correction

Faster postings merges:
Skip pointer

Recall basic merge

- Walk through the two postings simultaneously, in time linear in the total number of postings entries





If the list lengths are m and n , the merge takes $O(m+n)$ operations.

Can we do better?


Yes (if the index isn't changing too fast).

Augment postings with skip pointers (at indexing time)



- Why?
- To skip postings that will not figure in the search results.
- How?
- Where do we place skip pointers?

Query processing with skip pointers




Suppose we've stepped through the lists until we process **8** on each list. We match it and advance.

We then have **41** and **11** on the lower. **11** is smaller.

But the skip successor of **11** on the lower list is **31**, so we can skip ahead past the intervening postings.

Where do we place skips?

- Tradeoff:
 - More skips \rightarrow shorter skip spans \Rightarrow more likely to skip.
But lots of comparisons to skip pointers.
 - Fewer skips \rightarrow few pointer comparison, but then long skip spans \Rightarrow few successful skips.



Placing skips

- So... More skips or fewer skips... Where to add skip pointers???
- Simple heuristic: for postings of length L , use \sqrt{L} evenly-spaced skip pointers
- Easy if the index is relatively static; harder if L keeps changing because of updates.

Positional Index

Phrase queries

- Want to be able to answer queries such as “***stanford university***” – as a phrase
- Thus the sentence “*I went to university at Stanford*” is not a match.
 - The concept of phrase queries has proven easily understood by users; about 10% of web queries are phrase queries
- How??

A first attempt: Biword indexes

- Index every consecutive pair of terms in the text as a phrase
- For example the text “Friends, Romans, Countrymen” would generate the biwords
 - *friends romans*
 - *romans countrymen*
- Each of these biwords is now a dictionary term
- Two-word phrase query-processing is now immediate.

Longer phrase queries

- Longer phrases can be processed by breaking them down?
- ***stanford university palo alto*** can be broken into the Boolean query on biwords:

***stanford university AND university palo
AND palo alto***

Any problem?

Without the docs, we cannot verify that the docs matching the above Boolean query do contain the phrase.



Can have false positives!

Solution 2: Positional indexes

- In the postings, store, for each ***term*** the position(s) in which tokens of it appear:

<***term***, number of docs containing ***term***;

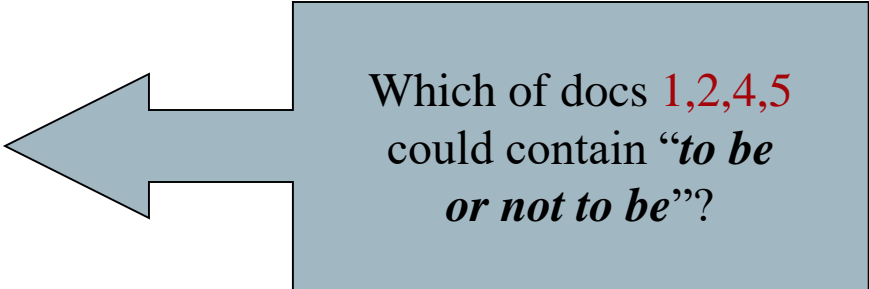
doc1: position1, position2 ... ;

doc2: position1, position2 ... ;

etc.>

Positional index example

<*be*: 993427;
1: 7, 18, 33, 72, 86, 231;
2: 3, 149;
4: 17, 191, 291, 430, 434;
5: 363, 367, ...>



Which of docs *1,2,4,5*
could contain “*to be*
or not to be”?

- Can compress position values/offsets
- Nevertheless, this expands postings storage *substantially*

Processing a phrase query

- Extract inverted index entries for each distinct term:
to, be, or, not.
- Merge their *doc:position* lists to enumerate all positions with “***to be or not to be***”.
 - ***to:***
 - 2:1,17,74,222,551; 4:8,16,190,429,433; 7:13,23,191; ...
 - ***be:***
 - 1:17,19; 4:17,191,291,430,434; 5:14,19,101; ...
- Same general method for proximity searches

Proximity queries

- LIMIT! /3 STATUTE /3 FEDERAL /2 TORT
 - Here, / k means “within k words of”.
- Clearly, positional indexes can be used for such queries; biword indexes cannot.

Positional index size

- Need an entry for each occurrence, not just once per document
- Index size depends on average document size
 - Average web page has <1000 terms
 - SEC filings, books, even some epic poems ... easily 100,000 terms
- Consider a term with frequency 0.1%

| Document size | Postings | Positional postings |
|---------------|----------|---------------------|
| 1000 | 1 | 1 |
| 100,000 | 1 | 100 |

Positional index size

- You can compress position values/offsets
- Nevertheless, a positional index expands postings storage substantially
- Nevertheless, it is now standardly used because of the power and usefulness of phrase and proximity queries ... whether used explicitly or implicitly in a ranking retrieval system.

Rules of thumb

- A positional index is 2–4 as large as a non-positional index
- Positional index size 35–50% of volume of original text

Positional Indexes: Wrap-up

- With a positional index, we can answer
 - phrase queries
 - proximity queries

Today...

- Need a better index than simple <term: docs>
- How can we improve on our basic index?
 - **Skip pointers**: faster postings merges
 - **Positional index**: Phrase queries and Proximity queries
 - **Permuterm index**: Wildcard queries